

Linear Algebra:

- Vector Multiplication: $x^T y = \sum_{i=1}^n x_i y_i$
- Matrix Multiplication: $C = AB \rightarrow C_{ij} = \sum_{k=1}^n A_{ik} B_{kj}$
- Gradients: $\nabla f = \left[\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right]^T$

- Gradients of Common Functions: $\nabla(a^T x) = a$, $\nabla(x^T Ax) = (A + A^T)x$

Machine Learning

- Supervised Learning: Given labelled data, attempt to learn function that assigns labels to unlabelled data
- Unsupervised Learning: Given unlabelled data, attempt to find structure
- K-nearest Neighbours: Classify based on majority label of k nearest points

Confusion matrix:	Predicted: yes	Actual: yes	Predicted: no	Actual: no
	TP	FN	FP	TN

- Precision = $\frac{TP}{TP+FP}$ Recall = $\frac{TP}{TP+FN}$ Specificity = $\frac{TN}{TN+FP}$

- Cross Validation: Split data into k subsets, train on k-1, test on 1, repeat k times

- Overfitting: Model fits training data too closely, fails to generalise

- Underfitting: Model too simple to capture underlying structure

- Decision Trees: Create best-split feature node, for all split data repeat

- Binary Classification: $f(x) \rightarrow y \in \{1, -1\}$

- Regression: $f(x) \rightarrow y \in \mathbb{R}$

- Hypothesis class: Function our algo will produce

- Loss function: How accurate is our predictor?

- Optimization Algo: How do we minimize loss?

Linear Regression:

- Hyp: $f(x) = w_1 + w_2 x$ or $f_w(x) = w \cdot \phi(x)$: $w = [w_1, w_2]$, $\phi(x) = [1, x]$
- Loss: $Loss(x, y, w) = (f_w(x) - y)^2$
- $TrainLoss(w) = \frac{1}{|D_{train}|} \sum_{(x,y) \in D_{train}} Loss(x, y, w)$
- Optimization: $\nabla_w TrainLoss(w) = \frac{1}{|D_{train}|} \sum_{(x,y) \in D_{train}} 2(w \cdot \phi(x) - y) \phi(x)$
- Gradient Update: $w \leftarrow w - \eta \nabla_w TrainLoss(w)$: $\eta = 0.1$

- Binary Classification: Create a linear function representing a boundary

- Hyp: $f(x) = sign(w \cdot \phi(x))$: $sign(z) = 1 : z > 0, -1 : z < 0, 0 : z = 0$
- Zero-One Loss: $Loss_{0-1}(x, y, w) = 1[f_w(x) \neq y]$
- Hinge Loss: $Loss_{hinge}(x, y, w) = \max\{1 - (w \cdot \phi(x))y, 0\}$
- Logistic Regression: $Loss_{logistic}(x, y, w) = \log(1 + e^{-(w \cdot \phi(x))y})$

- Stochastic Gradient Descent: Gradient descent is slow!

Neural Networks

- Hyp: $f(x) = V\sigma(W\phi(x))$ where σ is activation function
- Activation functions: $\sigma(z) = \max\{0, z\}$ (ReLU), $\sigma(z) = \frac{1}{1+e^{-z}}$ (sigmoid)
- Forward propagation: Compute layer outputs from input to output
- Backpropagation: Compute gradients using chain rule, update weights
- Architecture: Input layer \rightarrow Hidden layers \rightarrow Output layer
- Layers represent multiple layers of abstraction

- How to prevent overfitting?

- Reduce Dimensionality by removing features

- Regularize: $\min_w TrainLoss(w) + \lambda/2 \cdot ||w||^2$

Search:

No uncertainty

Definitions: b : branching factor, m : maximum depth

- Components of a search problem:

- State space: All possible states of an environment

- Successor function: function that maps actions to consequences

- Start state and end goal state

- DFS:
 - Idea: Expand one node as much as possible before exploring another path
 - $O(b^m)$ time complexity
 - $O(bm)$ space complexity
 - Complete if infinite tree size prevented and not optimal

BFS:

- Idea: prioritize nodes closest to start
- $O(b^d)$ time complexity where d is depth of solution
- $O(b^d)$ space complexity
- Complete and optimal if step costs are equal

Dijkstra's (Uniform Cost Search):

- Idea: prioritize expansion of low-cost paths
- $O(b^{1+\lfloor C^*/\epsilon \rfloor})$ time complexity where C^* is optimal cost, ϵ is minimum step cost
- $O(b^{1+\lfloor C^*/\epsilon \rfloor})$ space complexity
- Complete if step costs $\geq \epsilon > 0$ and optimal

Greedy Search

- Idea: Expand a node that we think is closest to the goal state

- Utilize a heuristic

A* Informed Search Algorithm

- Idea: use a heuristic to improve efficiency and make informed decisions
- Heuristic: A function that estimates how close a state is to the goal
- A heuristic is admissible if it never overestimates the cost of reaching a goal: $0 \leq h(n) \leq h^*(n)$ where $h(n)$ is the cost to the nearest goal
- A heuristic is consistent if $h(A) - h(C) \leq cost(A \rightarrow C)$. If it is consistent then it is optimal
- $f(n) = g(n) + h(n)$

Markov Decision Processes

- MDPs are defined by:

- set of states S and actions $a \in A$
- transition functions $T(s, a, s')$ and reward functions $R(s, a, s')$
- Start state and terminal state

- Utility: sum of rewards received

- Expected utility: $EU(a) = \sum_{a'} P(RES(a) = s')U(s')$

- Policies: the plan to get from start to goal: $\pi*: S \rightarrow A$

- With infinite utility we can mitigate with a finite horizon (limiting depth) or utilizing discounting: $U([(r_0, \dots, r_\infty)]) = \sum_{t=0}^\infty \gamma^t r_t \leq R_{max}/(1 - \gamma)$

- Discounting: Sooner rewards have higher utility than later rewards

- The value of a state s is: $V*(s) = \max_a Q*(s, a)$: $Q*(s, a)$ = expected utility starting out having taken action a from state s and acting optimally

- $\pi^*(s) = argmax_a Q*(s, a)$

- Bellman Equations: Take correct first action \rightarrow continue acting optimally

- $V*(s) = \max_a Q*(s, a)$

- $Q*(s, a) = \sum_{s'} T(s, a, s')[R(s, a, s') + \gamma V*(s')]$

- $V(s)$: Calculated value of a state

- $Q(s, a)$: Calculated value of an action from a state

- $\pi(s)$: What action we should take at a state

- Alternatively, we can perform policy iteration where we calculate values for some fixed policy until convergence, then update the policy using resulting converged values as future values

Reinforcement Learning

MDP's without T or R!

- MDPs, but we have to learn T and R
- Agent tries action A, gets $\{s, r\}$
- Model-Based Reinforcement Learning: Learn an approximate model based on experiences
 - STEP 1: Learn empirical MDP Model
 - Count $s' : \forall s, a$
 - Normalize to estimate $\hat{T}(s, a, s')$
 - Discover $\hat{R}(s, a, s')$ when we experience s, a, s'
 - STEP 2: Solve MDP, then run learned policy
- Model-Free Reinforcement Learning: Learn policies/values directly from experience
 - Passive RL: Execute a fixed policy $\pi(s)$ and learn state values
 - Direct Evaluation: Average sample values for each state
 - Temporal Difference Learning: Update $V^\pi(s)$ after each transition
 - TD Update: $V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + \alpha[R(s, \pi(s), s') + \gamma V^\pi(s')]$
 - Active RL: Learn optimal policy while acting
 - Q-Learning: Learn Q-values without knowing T or R
 - Q-Learning Update: $Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha[R(s, a, s') + \gamma \max_{a'} Q(s', a')]$
 - Exploration vs Exploitation: Balance trying new actions vs using known good actions
 - ϵ -greedy: With probability ϵ choose random action, else choose $argmax_a Q(s, a)$
 - We can approximate Q-learning with functions f_1, f_2, \dots
 - $V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$
 - $Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$

Games

- Types of games: Zero-sum: Competition (agents have opposite utilities), General-sum games: Agents have independent utilities
- We have states $s \in S$, players $p \in P$, actions $a \in A$, functions $S \times A \rightarrow S$, terminal test $S \rightarrow \{true, false\}$, terminal utilities: $S \times P \rightarrow R$
- Minimax: Time: $O(b^m)$, Space: $O(bm)$
- Alpha-Beta Pruning: Prune branches in a minimax tree to increase efficiency. Time: $O(b^{m/2})$.
- Multi-agent minimax: Terminals have utility tuples, node values are also utility tuples, each player maximizes its own component.
- Monte-Carlo Tree Search
 - Evaluation by rollouts: play multiple games to termination from a state s and count wins and losses
 - Selective search: explore parts of the tree that will help improve the decision at the root, regardless of depth
- Rollout: for each rollout, repeat until terminal: play a move according to a fixed, fast rollout policy. The fraction of wins correlates with the true value of the position
- MCTS V1: Allocate rollouts to promising and uncertain nodes.
- promising and uncertain are defined by: $UCB1(n) = \frac{U(n)}{N(n)} + C \times \sqrt{\frac{\log N(PARENT(n))}{N(n)}}$
 - $N(n)$ number of rollouts from node n
 - $U(n)$ total utility of rollouts for Player(Parent(n))
- Theorem: As $N \rightarrow \infty$ UCT selects the minimax move

Bayesian Networks

- $P(a|b) = \frac{P(a,b)}{P(b)}$
- $P(Y|e) = \frac{P(Y,e)}{P(e)} = \alpha P(Y,e) = \alpha \sum_{h \in H} P(Y, e, h) \rightarrow \alpha = \sum_{y \in Y} P(Y, e)$
- Normalization: multiply each entry by $\alpha = 1/(Sum \text{ of all entries})$
- Chain Rule: $P(x_1, x_2, \dots, x_n) = \prod_i P(x_i|x_1, \dots, x_{i-1})$
- Bayesian Statistics: Probabilities represent my uncertainty about the world
- Frequentist Statistics: Probabilities represent frequencies of real random outcomes. Properties of the world are fixed
- Probabilistic Inference: compute a desired probability from a probability model:
 - Enumerate options with sum
 - Sum out irrelevant variables
 - Normalize
- Time, Space, Data points: $O(d^n)$
- Bayes Rule: $p(H = h|Y = y) = \frac{P(H=h)P(Y=y|H=h)}{P(Y=y)}$, $P(a|b) = \frac{P(b|a)P(a)}{P(b)}$
- $P(Y = y) = \sum_{h' \in H} p(H = h')P(Y = y|H = h')$
- Two variables X and Y are independent if $\forall x, y : P(x, y) = P(x)P(y)$
- This implies that $P(x|y) = P(x)$, $P(y|x) = P(y)$
- Conditional Independence: X is conditionally independent of Y given Z iff:
 - $\forall x, y, z : P(x|y, z) = P(x|z)$ or $\forall x, y, z : P(x, y|z) = P(x|z)P(y|z)$
- Bayes' Nets: Complex joint distributions using conditional distributions
 - Nodes: Variables (with domains). Can be assigned (observed) or unassigned (unobserved)
 - Arcs: Interactions that indicate direct influence between variables
 - For a Bayesian net, we have a set of nodes: X_i for each variable, a directed acyclic graph, and a conditional distribution for each node given its parent variables in the graph
 - A conditional probability table is a table where each row is a distribution for the child given values of its parents
 - $P(X_1, \dots, X_n) = \prod_i P(X_i|Parents(X_i))$

Markov Models

- Markov Models have an initial distribution $P(x_0)$, transition model $P(X_t|X_{t-1})$
- Hidden Markov Models have an additional sensor model with observations $P(E_t|X_t)$
 - Filtering $P(X_t|e_{1:t})$
 - Prediction $P(X_{t+k}|e_{1:t})$
 - Smoothing: $P(X_k|e_{1:t})$: $k < t$
 - Explanation: $P(X_{1:t}|e_{1:t})$
- Markov Assumption: $P(X_t|X_0, \dots, X_{t-1}) = P(X_t|X_{t-1})$
- Stationary Assumption: Transition probabilities don't change over time
- Joint Distribution: $P(X_0, X_1, \dots, X_T) = P(X_0) \prod_{t=1}^T P(X_t|X_{t-1})$
- Filtering: $P(X_t|e_{1:t}) = \alpha P(e_t|X_t) \sum_{x_{t-1}} P(X_t|x_{t-1})P(x_{t-1}|e_{1:t-1})$
- (α represents the normalization operator here)
- Prediction: $P(X_{t+1}|e_{1:t}) = \sum_{x_t} P(X_{t+1}|x_t)P(x_t|e_{1:t})$
- Most Likely Explanation: $\arg\max_{x_{1:t}} P(x_{1:t}|e_{1:t})$ using Viterbi algorithm
- Viterbi: $m_{t+1}[x_{t+1}] = P(e_{t+1}|x_{t+1}) \max_{x_t} P(x_{t+1}|x_t)m_t[x_t]$
- Time steps: Elapse time (predict forward), observe (weight particles), resample

Algorithms:

- K-nearest Neighbours**
 function KNN(x, data, k):
 distances = []
 for each (x_i, y_i) in data:
 d = distance(x, x_i)
 distances.append(d, y_i))
 distances.sort()
 neighbors = distances[0:k]
 return mode([y for (d, y) in neighbors])
- Learning a Decision Tree**
 function LearnDecisionTree(data, attributes):
 if all examples have same label:
 return Leaf(label)
 if attributes is empty:
 return Leaf(majority_label(data))
 best_attr = choose_best_attribute(data, attributes)
 tree = Node(best_attr)
 for each value v of best_attr:
 subset = {x in data : x[best_attr] = v}
 if subset is empty:
 tree.add_branch(v, Leaf(majority_label(data)))
 else:
 remaining = attributes - {best_attr}
 subtree = LearnDecisionTree(subset, remaining)
 tree.add_branch(v, subtree)
 return tree
- function choose_best_attribute(data, attributes):
 return argmax over attributes of information_gain
- K-Means**
 function KMeans(data, k):
 centroids = randomly_select_k_points(data)
 repeat until convergence:
 clusters = [[] for i in range(k)]
 for each point x in data:
 closest = argmin_i distance(x, centroids[i])
 clusters[closest].append(x)
 for i in range(k):
 centroids[i] = mean(clusters[i])
 return centroids, clusters
- Stochastic Gradient Descent**
 function SGD(data, w_init, eta, num_epochs):
 w = w_init # [0,...,0]
 for epoch in range(num_epochs):
 shuffle(data)
 for each (x, y) in data:
 gradient = compute_gradient(Loss(x, y, w), w)
 w = w - eta * gradient
 return w
- DFS**
 function DFS(problem):
 frontier = Stack()
 frontier.push(start_state)
 explored = set()
 while frontier is not empty:
 node = frontier.pop()
 if node is goal:
 return solution
 explored.add(node)
 for each successor of node:
 if successor not in explored and not in frontier:
 frontier.push(successor)
 return failure
- BFS**
 function BFS(problem):
 frontier = Queue()
 frontier.enqueue(start_state)
 explored = set()
 while frontier is not empty:
 node = frontier.dequeue()
 if node is goal:
 return solution
 explored.add(node)
 for each successor of node:
 if successor not in explored and not in frontier:
 frontier.enqueue(successor)
 return failure
- Dijkstra's Uniform Cost Search Algorithm**
 function Dijkstra(problem):
 frontier = PriorityQueue()
 frontier.push(start_state, 0)
 explored = set()
 cost = {start_state: 0}
 while frontier is not empty:
 node = frontier.pop()
 if node is goal:
 return solution
 explored.add(node)
 for each successor of node:
 new_cost = cost[node] + step_cost(node, successor)
 if successor not in explored:
 if successor not in cost or new_cost < cost[successor]:
 cost[successor] = new_cost
 frontier.push(successor, new_cost)
 return failure
- A* Algorithm**
 function AStar(problem):
 frontier = PriorityQueue()
 frontier.push(start_state, h(start_state))
 explored = set()
 g_cost = {start_state: 0}
 while frontier is not empty:
 node = frontier.pop()
 if node is goal:
 return solution
 explored.add(node)
 for each successor of node:
 new_cost = g_cost[node] + step_cost(node, successor)
 if successor not in explored:
 if successor not in g_cost or new_cost < g_cost[successor]:
 g_cost[successor] = new_cost
 f_cost = new_cost + h(successor)
 frontier.push(successor, f_cost)
 return failure
- Expectimax:** For games with chance nodes, compute expected values instead of min/max
 function EM(state):
 if state is terminal:
 return utility(state)
 if state is max node:
 return max over actions of EM(successor(state, action))
 if state is chance node:
 return sum outcomes of P(outcome) * EM(successor(state, outcome))
- Minimax**
 function Minimax(state):
 if state is terminal:
 return utility(state)
 if state is max node:
 return max over actions of Minimax(successor(state, action))
 if state is min node:
 return min over actions of Minimax(successor(state, action))
- Alpha Beta Pruning**
 function ABP(state, alpha, beta):
 if state is terminal:
 return utility(state)
 if state is max node:
 v = -infinity
 for each action:
 v = max(v, ABP(successor(state, action), alpha, beta))
 if v >= beta:
 return v # Beta cutoff
 alpha = max(alpha, v)
 return v
 if state is min node:
 v = +infinity
 for each action:
 v = min(v, ABP(successor(state, action), alpha, beta))
 if v <= alpha:
 return v # Alpha cutoff
 beta = min(beta, v)
 return v
- MCTS Version 2.0 UCT**
 function MCTS(root_state):
 tree = {root_state: {}}
 N = {root_state: 0}
 U = {root_state: 0}
 repeat until out of time:
 # Selection: traverse tree using UCB1
 node = root_state
 path = [node]
 while node is fully expanded and not terminal:
 node = argmax_child_UCB1(child)
 path.append(node)
 # Expansion: add new child
 if node is not terminal:
 child = unexplored_child(node)
 tree.add(child)
 N[child] = 0
 U[child] = 0
 path.append(child)
 # Simulation: rollout from new node
 result = rollout(child)
 # Backpropagation: update counts
 for n in path:
 N[n] += 1
 U[n] += result
 # Return action with highest visit count
 return argmax_child N[child]
- Value Iteration**
 function ValueIteration(S, A, T, R, gamma):
 Initialize V(s) = 0 for all s
 repeat until convergence:
 V_new = {}
 for each state s in S:
 V_new[s] = max_a sum_s' T(s,a,s')[R(s,a,s') + gamma*V(s')]
 V = V_new
 return V, extract_policy(V)
- Policy Iteration**
 function PolicyIteration(S, A, T, R, gamma):
 Initialize random policy pi
 repeat until policy stable:
 # Policy Evaluation
 V = solve V(s) = sum_s' T(s,pi(s),s')[R(s,pi(s),s') + gamma*V(s')]
 # Policy Improvement
 pi_new(s) = argmax_a sum_s' T(s,a,s')[R(s,a,s') + gamma*V(s')]
 if pi_new == pi: break
 pi = pi_new
 return pi
- Forward Algorithm for Hidden Markov Models**
 function Forward(observations, T, E, pi):
 alpha_0 = pi * E[obs_0]
 for t = 1 to T:
 alpha_t = E[obs_t] * (T^t * alpha_{t-1})
 normalize(alpha_t)
 return alpha_T
- Backward Algorithm**
 function Backward(observations, T, E):
 beta_T = [1, 1, ..., 1]
 for t = T-1 down to 0:
 beta_t = T * (E[obs_{t+1}] * beta_{t+1})
 normalize(beta_t)
 return beta_0
- Equations**
 - Exponential Moving Average: $x_n = \alpha x_{n-1} + (1 - \alpha)x_n$
 - Log-Based Normalization: $\frac{1}{T} \log(P(w_0, \dots, w_T)) = \frac{1}{T} \sum_{t=1}^T \log(P(w_t | w_{t-1}))$
 - HMM Forward Pass: $\alpha_t(j) = P(O_t | X_t = j) \sum_i \alpha_{t-1}(i) \cdot P(X_t = j | X_{t-1} = i)$
 - Entropy: $H(X) = -\sum_i P(x_i) \log_2 P(x_i)$
 - Information Gain: $IG(Y|X) = H(Y) - \sum_{x \in X} P(x)H(Y|X = x)$
 - Learning rate decay: $\alpha_t = \frac{1}{1+t}$
 - L2 Regularization: $\lambda ||w||_2^2 = \lambda \sum_i w_i^2$
 - Sum of geometric series: $\sum_{t=0}^{\infty} \gamma^t = \frac{1}{1-\gamma}$ for $|\gamma| < 1$
 - Conditional probability expansion: $P(A, B|C) = P(A|B, C)P(B|C)$
 - Law of Total Probability: $P(A) = \sum_i P(A|B_i)P(B_i)$
 - Marginalization: $P(X) = \sum_y P(X, Y = y)$
 - Product Rule: $P(X, Y) = P(X|Y)P(Y) = P(Y|X)P(X)$
 - Total Probability: $P(Y) = \sum_x P(Y|X = x)P(X = x)$
 - Independence: $P(X, Y) = P(X)P(Y) \iff P(X|Y) = P(X)$
 - Conditional Independence: $P(X, Y|Z) = P(X|Z)P(Y|Z)$
 - Complete Bellman Equation: $V^*(s) = \max_a \sum_{s'} T(s, a, s')[R(s, a, s') + \gamma V^*(s')]$
 - Policy Extraction via Q-Values: $\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$
 - Q-Value Iteration: $Q_{k+1}(s, a) = \sum_{s'} T(s, a, s')[R(s, a, s') + \gamma \max_{a'} Q_k(s', a')]$